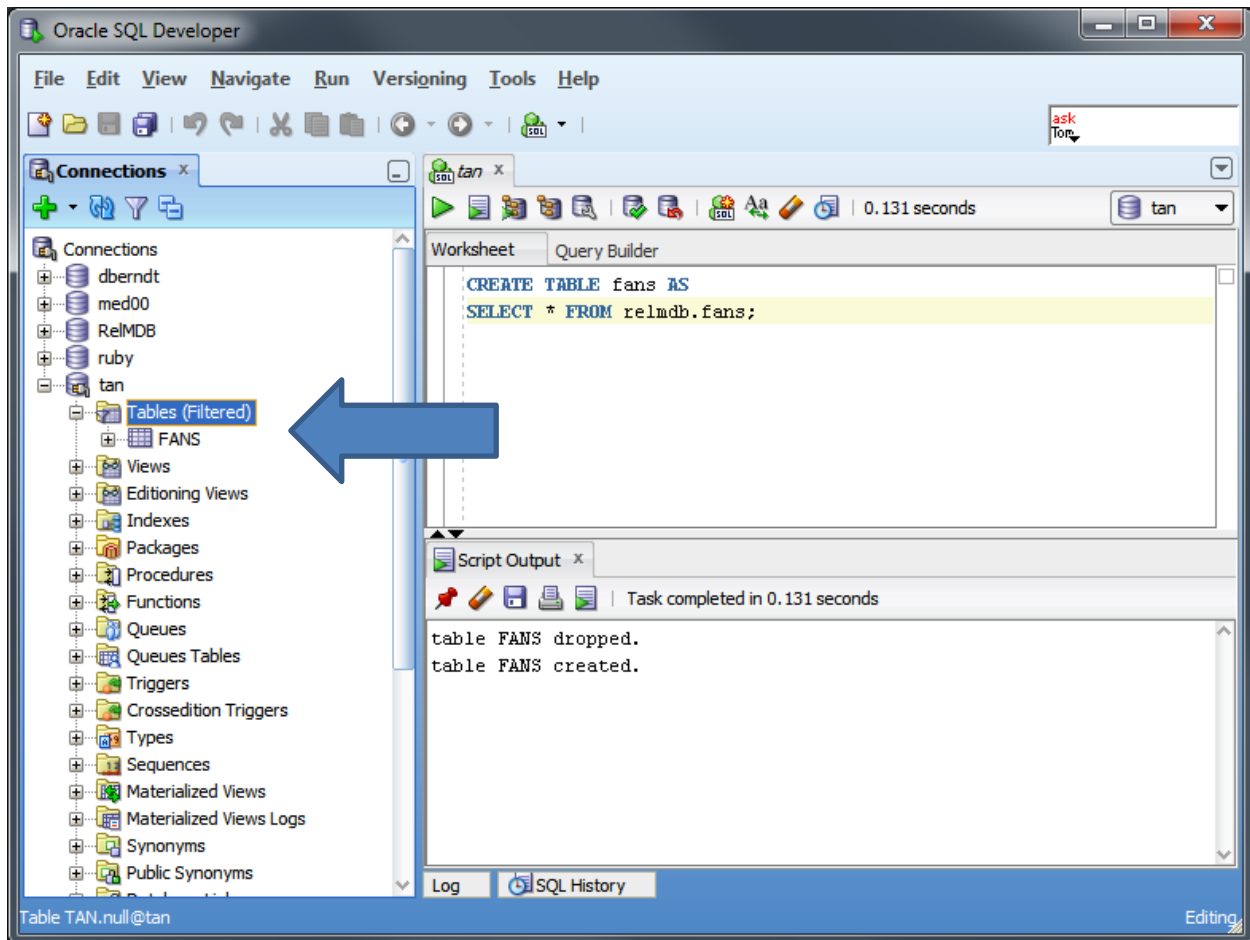


## Assignment 1 “Cheat Sheet” for SQL Query Writing

### Getting Started

This document is intended to help you complete the query writing assignment, especially if you are rusty or have not seen a lot of SQL before. Though, it also includes some basic features of the SQL Developer environment which all might find helpful.

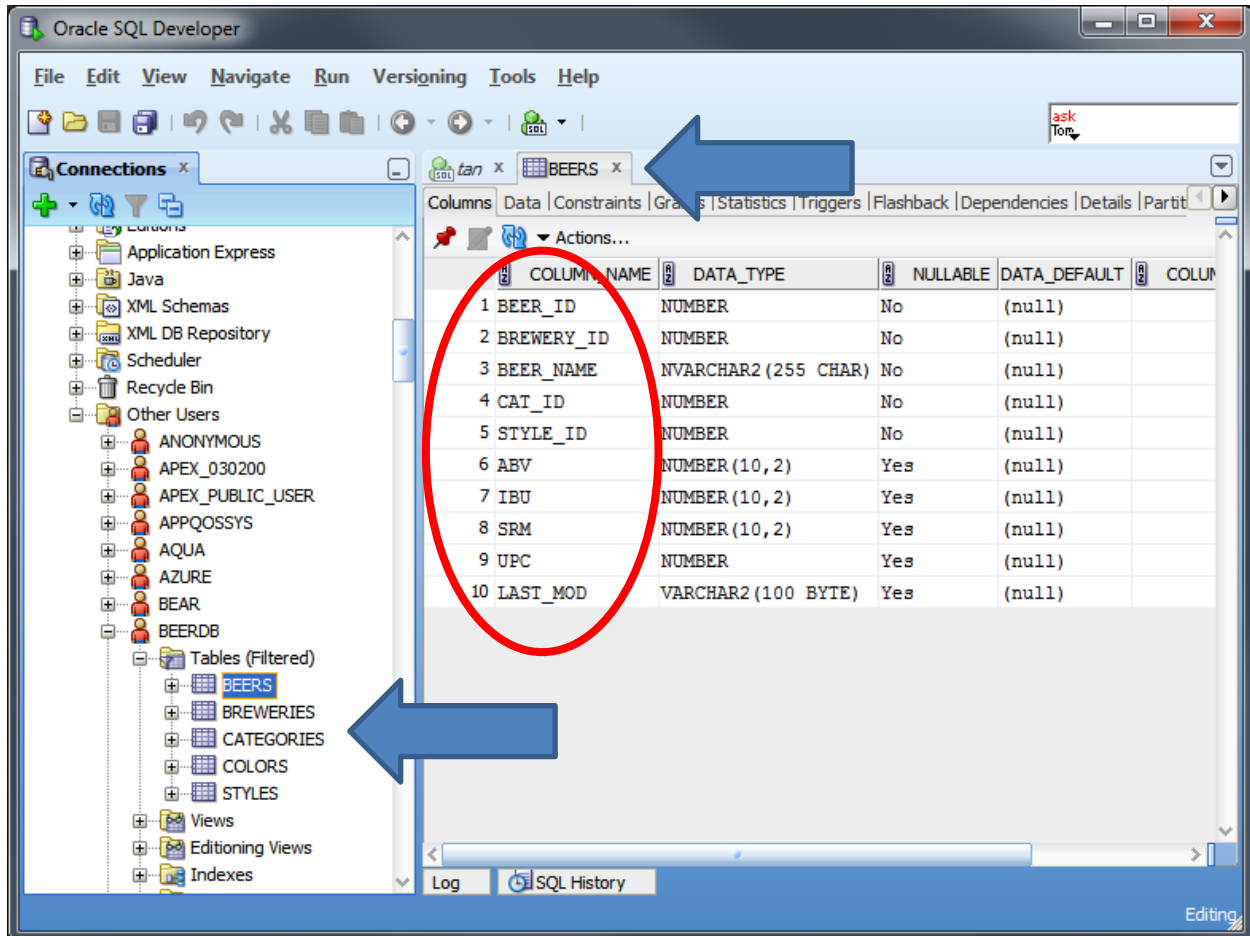
First make a connection to the database using your user login (here I used TAN). All the folders you see below the user name contain database objects you own. So, these folders should be empty if you are using a fresh account. The CREATE TABLE AS SELECT (CTAS) statement shown here references the FANS table in the RELMDB schema and creates a copy here. Note: There is a FANS table in the folder after refreshing.



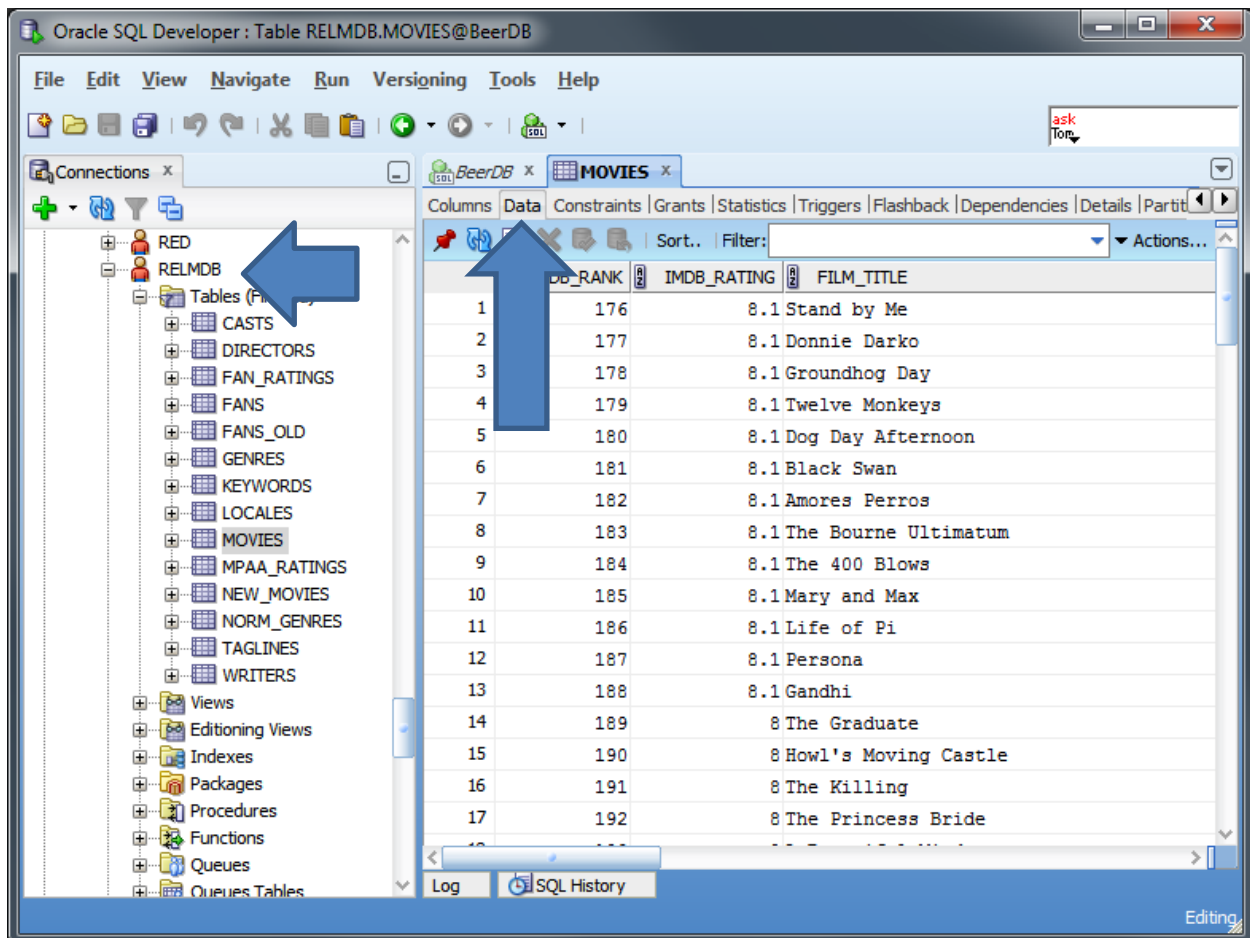
We really don't need any tables of our own yet, since we can write queries against the movie database (RELMDB) and craft beer database (BEERDB).

## Where are the beer and movie databases?

The Oracle database instance includes many users and associated schemas. So, look for the “Other Users” folder and expand the list. You can browse any schema, but look specifically for the BEERDB and RELMDB users and expand those to see the tables. These are the tables you can query against by prefacing them with the schema name (such as BEERDB.BEERS). You can click on the table name to see the columns and other details about the table, as well as expand the table to see a column list. In this way, you can see what is available for query writing.



So, you can do the same with the RELMDB user to see the movie database. Again you can click on a table, such as the MOVIES table to see the details, including the actual data. This is a great way to get familiar with the data before writing any queries.



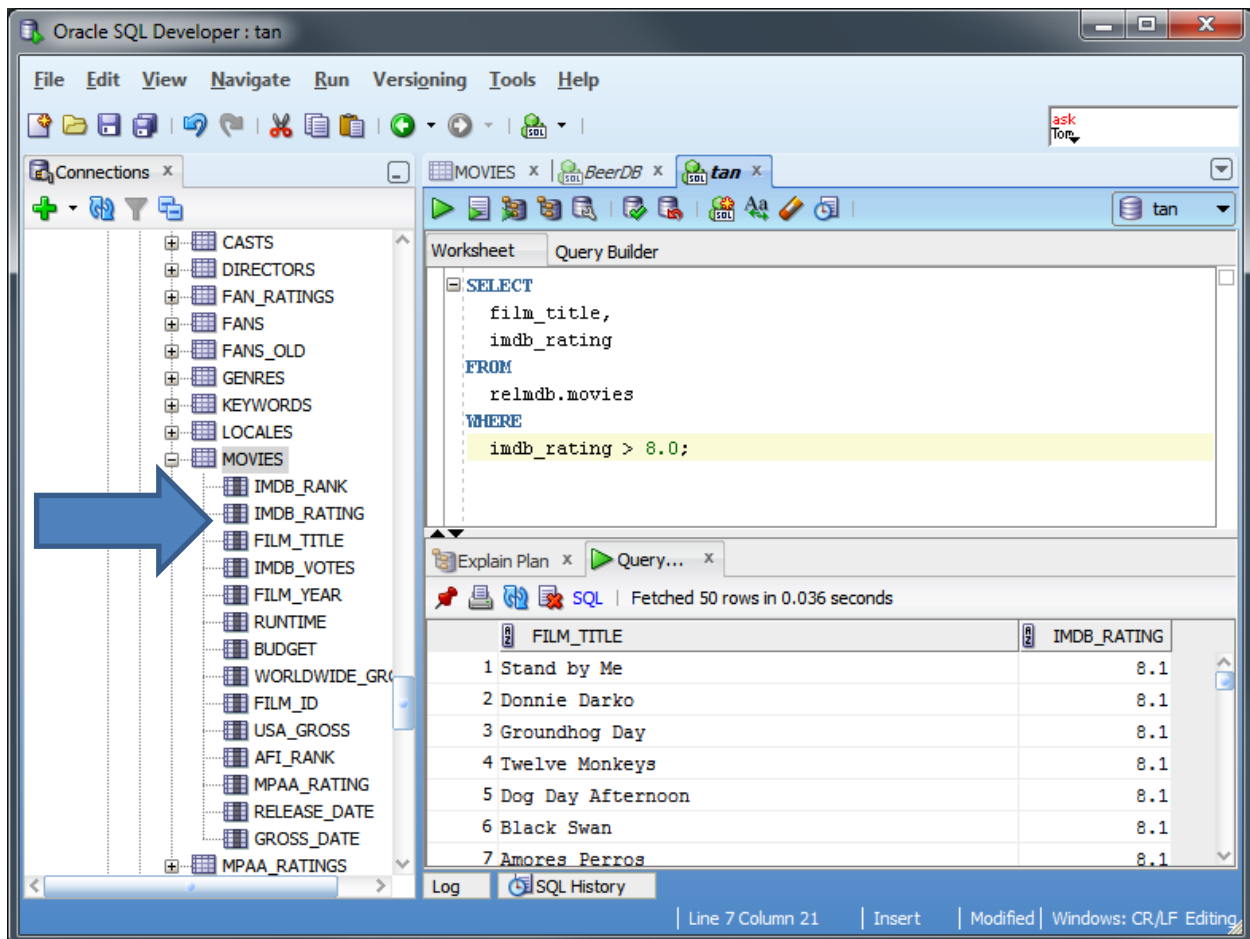
So, we are ready to write our first query for the assignment. Remember that as you type queries in the SQL worksheet, the tool is constantly checking the syntax (highlighting correct SQL keywords) and underlining problems. In addition, there will be suggestions for auto-completing parts of the queries.

## Writing Our First Query

So, we can start with the first query against the movie database. The description of the query is as follows.

*“Display the name of all movies that have an IMDB rating of at least 8.0, with more than 100,000 IMDB votes, and were released from 2007 to 2013. Show the movies with the highest IMDB ratings first.”*

The best way to tackle SQL query writing is in steps, building and testing simpler components and then putting the whole query together. You can execute the query along the way to see if the results meet your expectations. So, we can start by building the basic SELECT-FROM-WHERE clause structure. The target of our query is the MOVIES table, so that is referenced in the FROM clause as RELMDB.MOVIES. We will want to see the FILM\_TITLE and the IMDB\_RATING since these are integral to the query. Next, add the WHERE clause to find only those movies with IMDB\_RATING > 8.0. Go ahead and run the query and make sure that all is well. You can see the film titles and note that all the IMDB\_RATING values are above 8.0; so far so good.



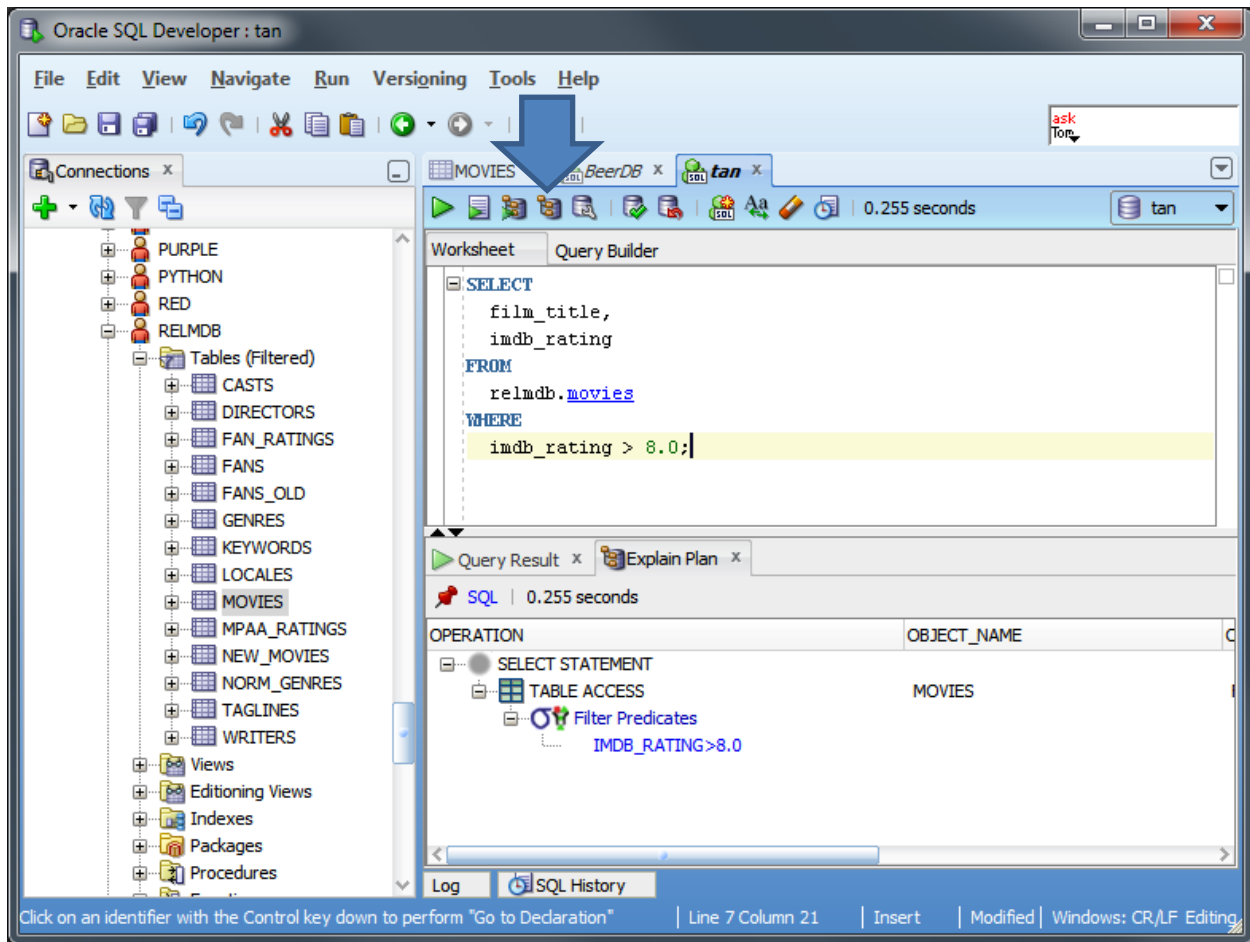
The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows a tree of database objects, with 'MOVIES' selected. A blue arrow points to the 'MOVIES' table. The 'Query Builder' pane shows the following SQL query:

```
SELECT
  film_title,
  imdb_rating
FROM
  relmdb.movies
WHERE
  imdb_rating > 8.0;
```

The 'Query' pane shows the results of the query, fetched in 0.036 seconds. The results are displayed in a table with two columns: 'FILM\_TITLE' and 'IMDB\_RATING'.

FILM_TITLE	IMDB_RATING
1 Stand by Me	8.1
2 Donnie Darko	8.1
3 Groundhog Day	8.1
4 Twelve Monkeys	8.1
5 Dog Day Afternoon	8.1
6 Black Swan	8.1
7 Amores Perros	8.1

Remember that SQL is a declarative language. You are really just describing the result you would like to see. You are asking to see the `FILM_TITLE` and `IMDB_RATING` attributes for just those movies with an `IMDB_RATING` above 8.0. The database query optimizer figures out how to compute the desired result set using a reasonably efficient execution plan. The “Explain Plan” button will show you exactly how your query is being executed. Here, the execution plan notes that it is a `SELECT` and that it is accessing the table `MOVIES` (with a `FULL` scan). It also includes the filter predicate `IMDB_RATING > 8.0`. We will look at execution plans in more depth as the course progresses.



We can continue building the query by adding the second WHERE predicate. We want movies that also have over IMDB\_VOTES > 100000. Note that we are capitalizing SQL keywords and writing table and column names in lower case. That is just a convention and many database programmers have other favorite style elements. Indentation and parentheses also help make the query easier to read. All the capitalization and whitespace is ignored by the database engine. Now our WHERE clause includes two predicates with AND requiring that both evaluate to TRUE. Go ahead and run the query again. The result set should be smaller since we are being more restrictive.

The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows a connection to 'tan' with a tree view of the 'RELIMDB' schema, including tables like CASTS, DIRECTORS, FAN\_RATINGS, FANS, FANS\_OLD, GENRES, KEYWORDS, LOCALES, and MOVIES. The 'MOVIES' table is selected, showing columns like IMDB\_RANK, IMDB\_RATING, FILM\_TITLE, IMDB\_VOTES, FILM\_YEAR, RUNTIME, BUDGET, WORLDWIDE\_GROSS, FILM\_ID, USA\_GROSS, AFI\_RANK, and MPAA\_RATING.

The 'Query Builder' pane shows the following SQL query:

```
SELECT
    film_title,
    imdb_rating
FROM
    relimdb.movies
WHERE
    (imdb_rating > 8.0) AND
    (imdb_votes > 100000);
```

The 'Script Output' pane shows the results of the query, with 135 rows fetched in 0.115 seconds. The results are displayed in a table with two columns: 'FILM\_TITLE' and 'IMDB\_RATING'.

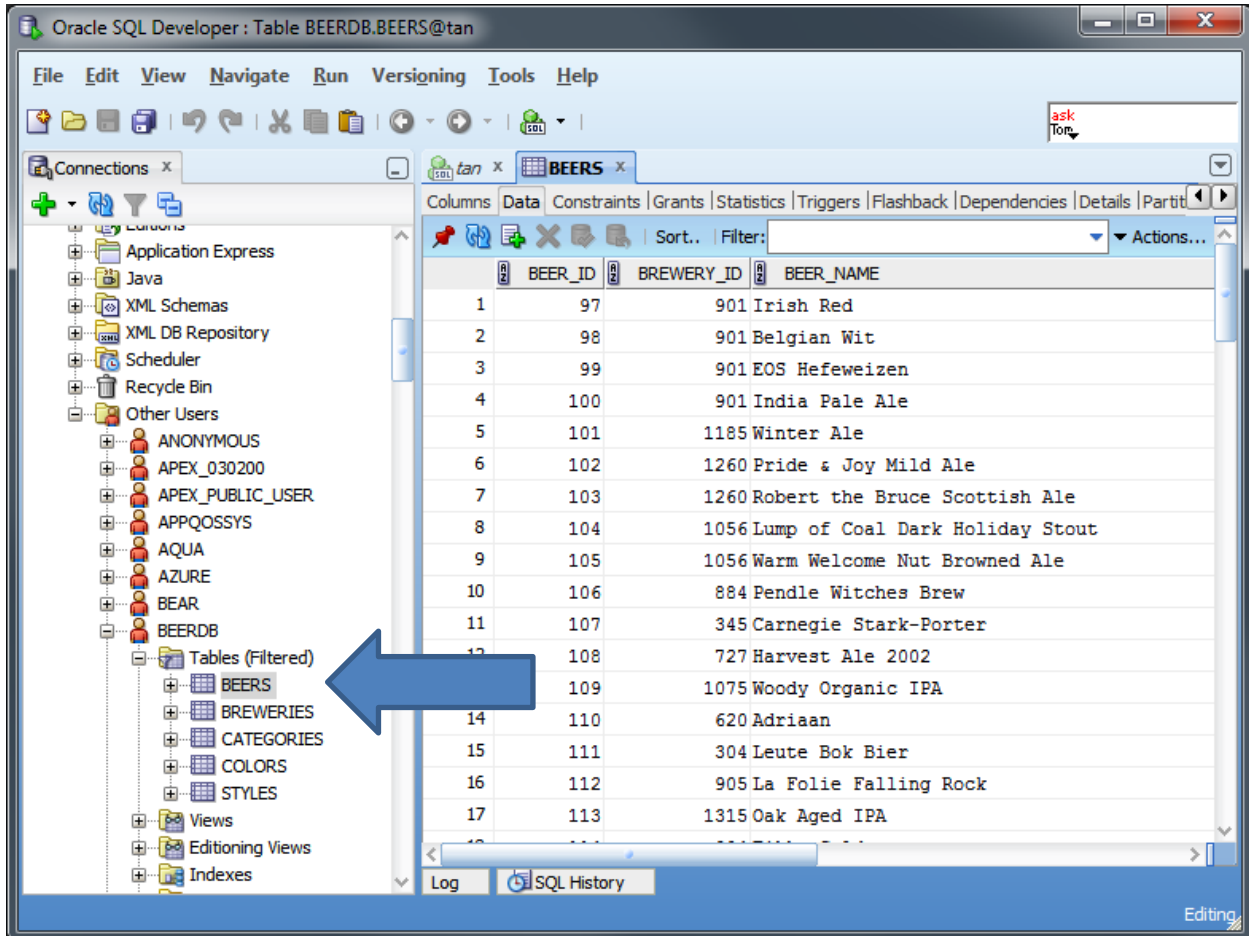
FILM_TITLE	IMDB_RATING
1 Stand by Me	8.1
2 Donnie Darko	8.1
3 Groundhog Day	8.1
4 Twelve Monkeys	8.1
5 Dog Day Afternoon	8.1
6 Black Swan	8.1
7 Amores Perros	8.1

The status bar at the bottom indicates 'Table RELIMDB.MOVIES@tan', 'Line 8 Column 25', 'Insert', 'Modified', and 'Windows: CR/LF Editing'.

To finish up, the query also asks for those movies that were released between 2007 and 2013. How should we add the additional predicates to the WHERE clause? You could add something like (FILM\_YEAR >= 2007) AND (FILM\_YEAR <= 2013) or look at the BETWEEN operator that seems like a perfect fit. That last specification is to show the "highest IMDB ratings first," which means we need to sort the result set. The ORDER BY clause provides that functionality. We will need to add ORDER BY IMDB\_RATING DESC to the query, yielding a SELECT-FROM-WHERE-ORDER BY structure. The sort order can be specified as ascending (ASC - the default) or descending (DESC). Here we need to use DESC to see the highest values first. Go ahead and complete the query.

## Writing Our Second Query

Let's try to write a second query, this time targeting the beer database (BEERDB). Again, looking at the main BEERS table, we can get a sense of the data. We can see that there is a BREWERY\_ID column that would be useful for joining with the BREWERIES table. So, we can start here.



Oracle SQL Developer: Table BEERDB.BEERS@tan

File Edit View Navigate Run Versigning Tools Help

Connections x

tan x BEERS x

Columns Data Constraints Grants Statistics Triggers Flashback Dependencies Details Partit

Sort.. Filter: Actions...

BEER_ID	BREWERY_ID	BEER_NAME
1	97	901 Irish Red
2	98	901 Belgian Wit
3	99	901 EOS Hefeweizen
4	100	901 India Pale Ale
5	101	1185 Winter Ale
6	102	1260 Pride & Joy Mild Ale
7	103	1260 Robert the Bruce Scottish Ale
8	104	1056 Lump of Coal Dark Holiday Stout
9	105	1056 Warm Welcome Nut Browned Ale
10	106	884 Pendle Witches Brew
11	107	345 Carnegie Stark-Porter
12	108	727 Harvest Ale 2002
13	109	1075 Woody Organic IPA
14	110	620 Adriaan
15	111	304 Leute Bok Bier
16	112	905 La Folie Falling Rock
17	113	1315 Oak Aged IPA

Log SQL History

Editing

First, we can try a simple single table query against the BREWERIES table. Find all the breweries in the state of Colorado. Later, we can join to the BEERS table to find out how many beers each brewery produces. So the query would be as follow (feel free to add this one to the assignment).

*“List all the breweries in Colorado, along with the beers they produce. Sort the results by brewery, so the beers appear together in the result set.”*

Here is the SQL for the first step, along with the result set. Our goal is to simply list the breweries in Colorado. On the left, we can see the available columns for the BREWERIES table, including STATE. There are definitely some interesting breweries.

The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows a tree view of the database schema. The 'BEERDB' connection is selected, and the 'Tables (Filtered)' list shows the 'BREWERIES' table. A blue arrow points from the 'STATE' column in the 'BREWERIES' table to the 'WHERE' clause of the SQL query in the 'Query Builder' pane. The query is:

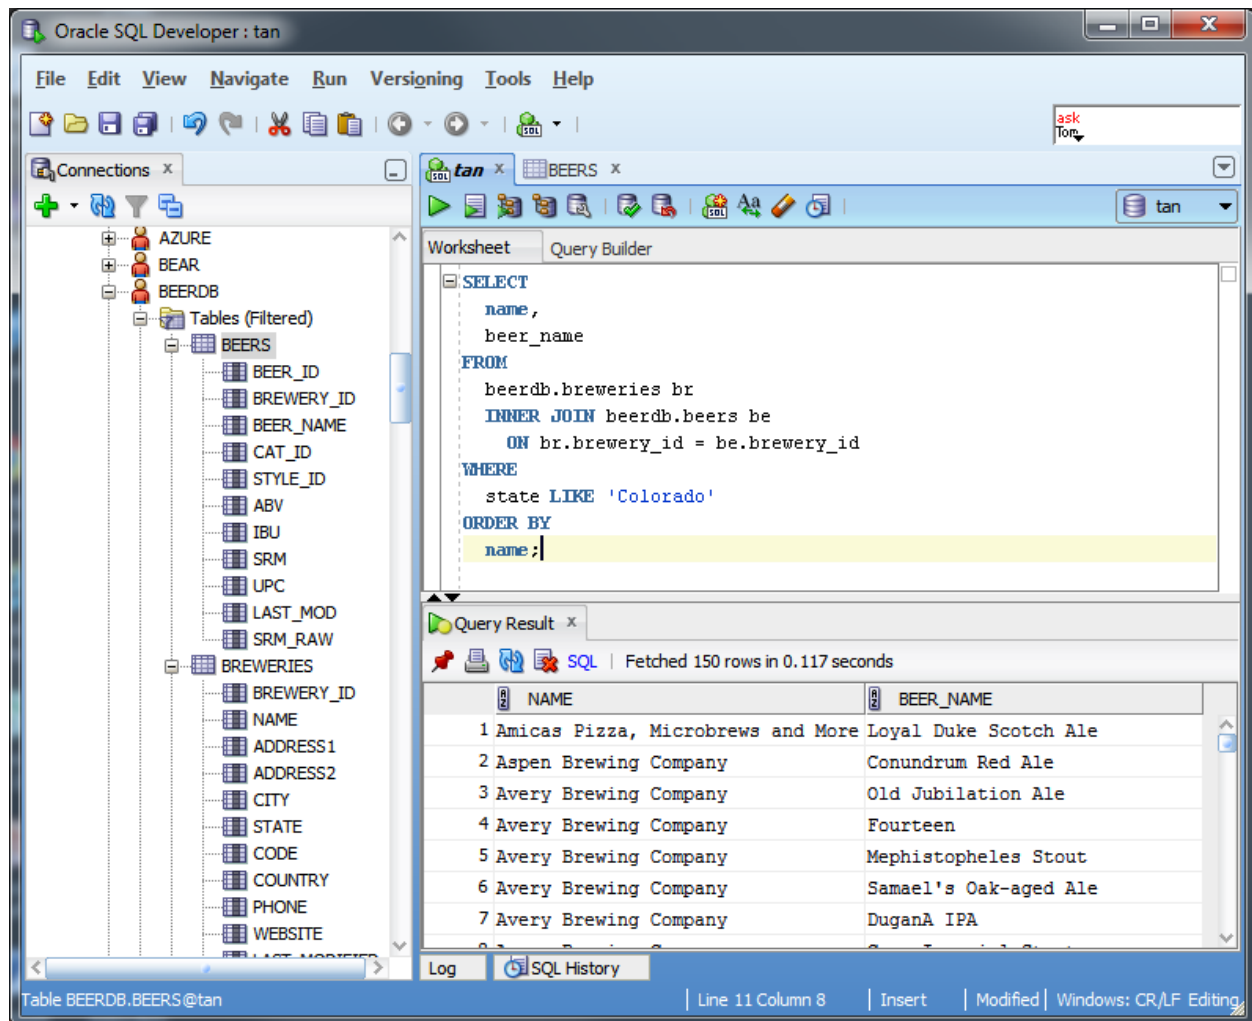
```
SELECT
  name,
  state
FROM
  beerdb.breweries
WHERE
  state LIKE 'Colorado'
```

The 'Query Result' pane shows the results of the query, fetched in 0.037 seconds. The results are displayed in a table with two columns: 'NAME' and 'STATE'.

NAME	STATE
1 Hops Grillhouse & Brewery - Cherry Creek	Colorado
2 Hubcap Brewery and Kitchen	Colorado
3 Jarre Creek Ranch Brewing	Colorado
4 Left Hand Brewing Company	Colorado
5 Amicas Pizza, Microbrews and More	Colorado
6 Avery Brewing Company	Colorado
7 Backcountry Brewery	Colorado

Again, we are building our query in steps. So, next we will need to join to the BEERS table to find all the beers produced. The BREWERY\_ID column is the obvious choice for joining the two tables. It also makes sense to use a basic INNER JOIN, the most widely used method for combining tables.





So, the query now uses an INNER JOIN to combine BREWERIES and BEERS. Notice that table aliases (“br” and “be”) are used to save some typing in the ON clause. This is called explicit join notation since INNER JOIN is specifically stated along with the join criteria in the ON clause. Nothing about the join is put into the WHERE clause. This is how you should write your joins.

Another interesting query could be built from here (please feel free to include this one in your assignment).

*“List all the breweries in Colorado, along with the number of beers they produce. Sort the results in descending order by the number of beers, so the most prolific breweries appear first.”*

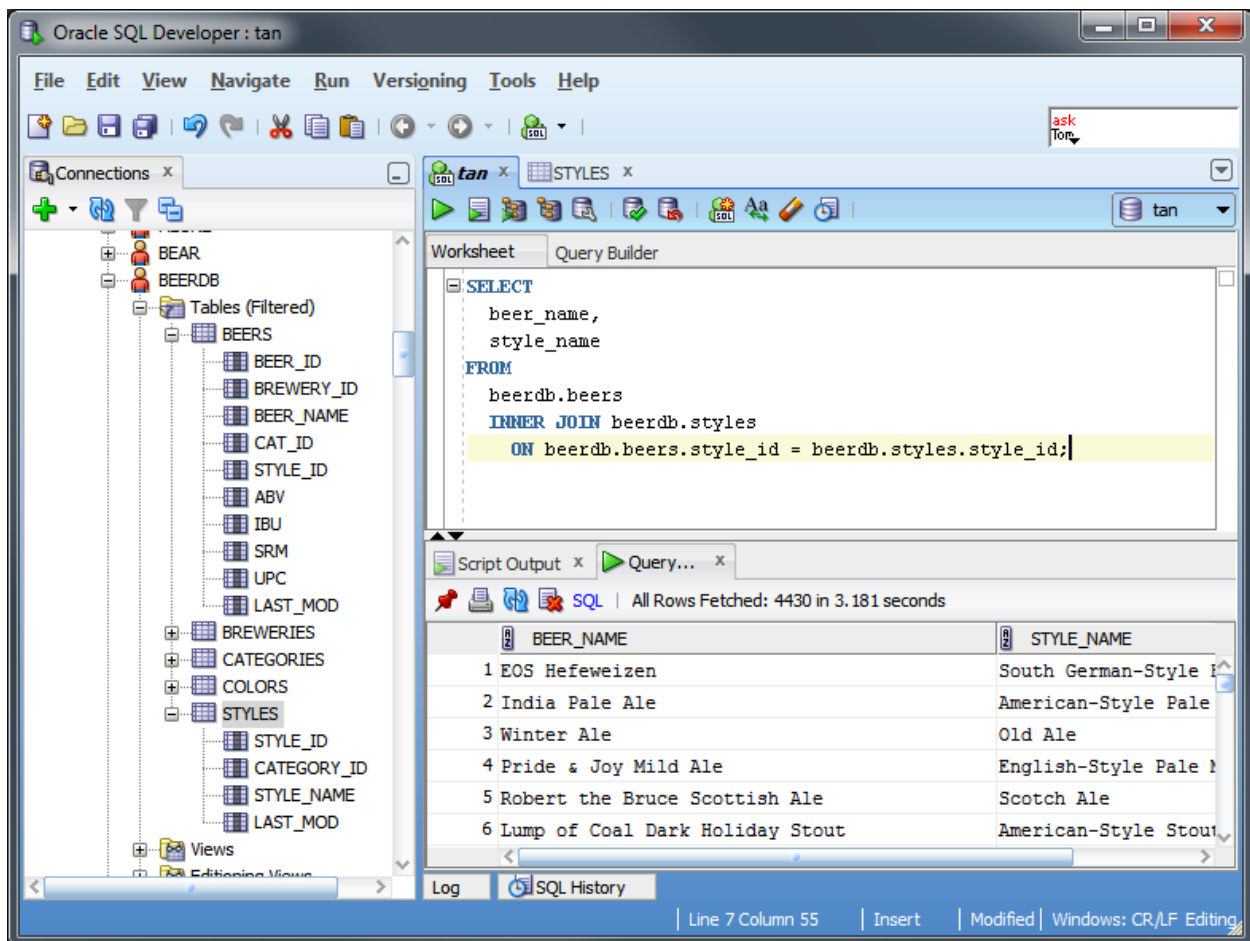
To write this query, you will need to replace the beer names with COUNT(beer\_name) to get the number (not a list of the beers). An aggregation operator like COUNT is typically used with a GROUP BY clause. Here you would want to GROUP BY the brewery name. Try to complete this query and include it in your assignment!

## Writing One More Beer Database Query

We can start on another query from the beer database to illustrate some more query writing skills. For example, the following query takes us further into different join types.

*“Display each beer’s name and style name. A beer should be displayed regardless of whether a style name exists or not.”*

We could start by simply joining BEERS with STYLES following the form outlined in the previous query. The query is fine and produces a reasonable answer set. However, the specification states that we should see a beer even if a style name does not exist. How do we handle this requirement?



The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows the 'BEERDB' database. The 'Tables (Filtered)' list includes BEERS, BEER\_ID, BREWERY\_ID, BEER\_NAME, CAT\_ID, STYLE\_ID, ABV, IBU, SRM, UPC, LAST\_MOD, BREWERIES, CATEGORIES, COLORS, and STYLES. The 'STYLES' table is selected. In the center, the 'Query Builder' window shows the following SQL query:

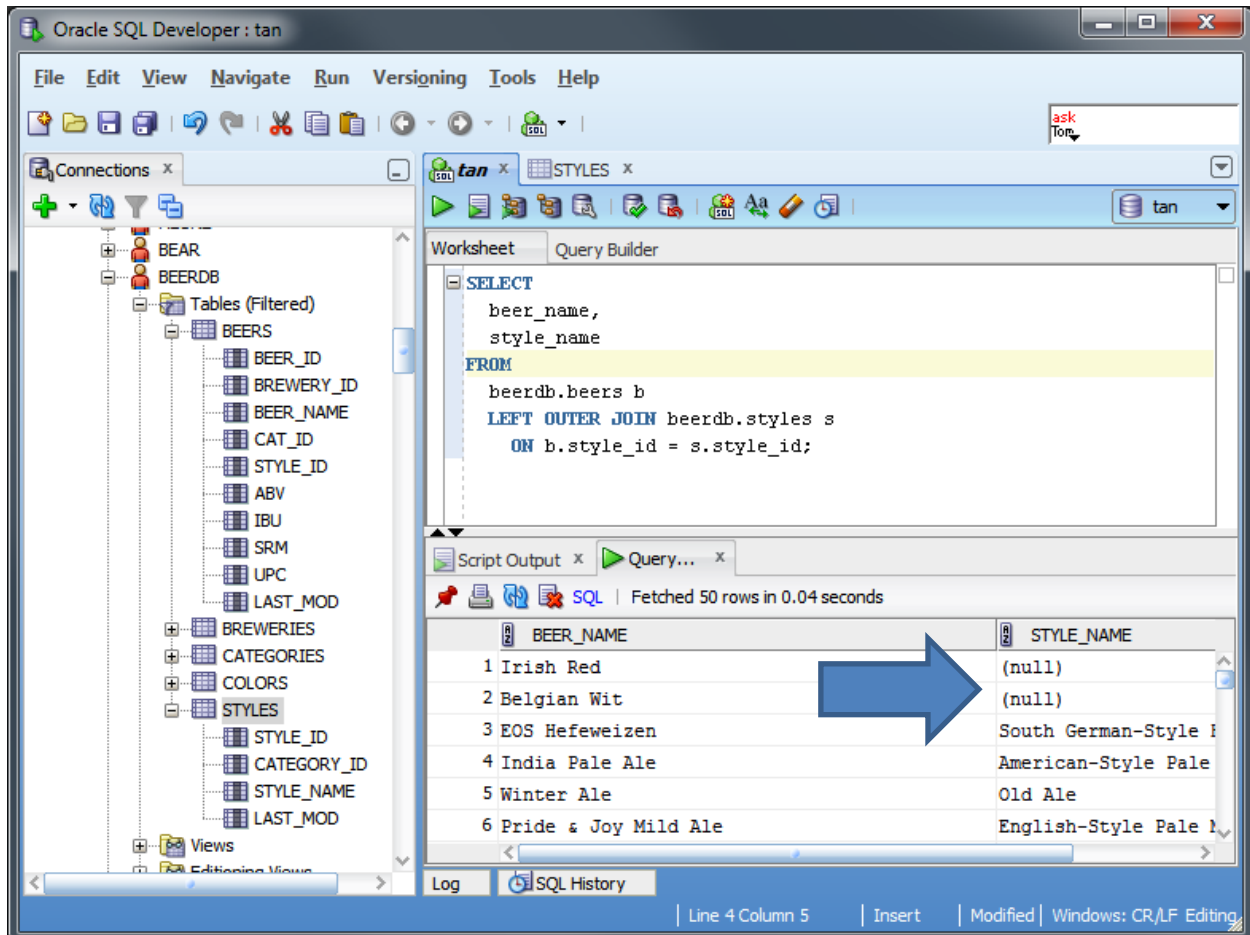
```
SELECT
  beer_name,
  style_name
FROM
  beerdb.beers
  INNER JOIN beerdb.styles
    ON beerdb.beers.style_id = beerdb.styles.style_id;
```

Below the query, the 'Script Output' window shows the results of the query. It indicates 'All Rows Fetched: 4430 in 3.181 seconds'. The results are displayed in a table with two columns: BEER\_NAME and STYLE\_NAME.

BEER_NAME	STYLE_NAME
1 EOS Hefeweizen	South German-Style P
2 India Pale Ale	American-Style Pale
3 Winter Ale	Old Ale
4 Pride & Joy Mild Ale	English-Style Pale M
5 Robert the Bruce Scottish Ale	Scotch Ale
6 Lump of Coal Dark Holiday Stout	American-Style Stout

There are different join types. The most commonly used type is the INNER JOIN, which means that the condition must be met in both tables. Here this means that there must be a STYLE\_ID in the beers table and a corresponding entry in the STYLES table. The OUTER JOIN is used when you want to include entries from one table even when not paired with an entry from the other table. A LEFT OUTER JOIN (or just LEFT JOIN) includes rows from the left-hand side even if not paired with an entry from the other table. A RIGHT OUTER JOIN just comes at it from the other side. A FULL OUTER JOIN includes rows from either side, even if there is no paired row from the other table. So, which one do we want here?

We want to use the LEFT OUTER JOIN as shown below, so that all the beer names from the BEERS table (the left-hand side) show in the query results. You can see here that some of the style names have NULL values. These are the rows that contain a beer name, but no corresponding style name. This version meets the query specification to show beer names “regardless of whether a style name exists or not.”



The screenshot displays the Oracle SQL Developer interface. On the left, the 'Connections' pane shows the 'BEER' database structure, including tables like BEERS, BEER\_ID, BREWERY\_ID, BEER\_NAME, CAT\_ID, STYLE\_ID, ABV, IBU, SRM, UPC, LAST\_MOD, BREWERIES, CATEGORIES, COLORS, and STYLES. The 'STYLES' table is selected. The 'Worksheet' pane shows the following SQL query:

```
SELECT
  beer_name,
  style_name
FROM
  beerdb.beers b
  LEFT OUTER JOIN beerdb.styles s
    ON b.style_id = s.style_id;
```

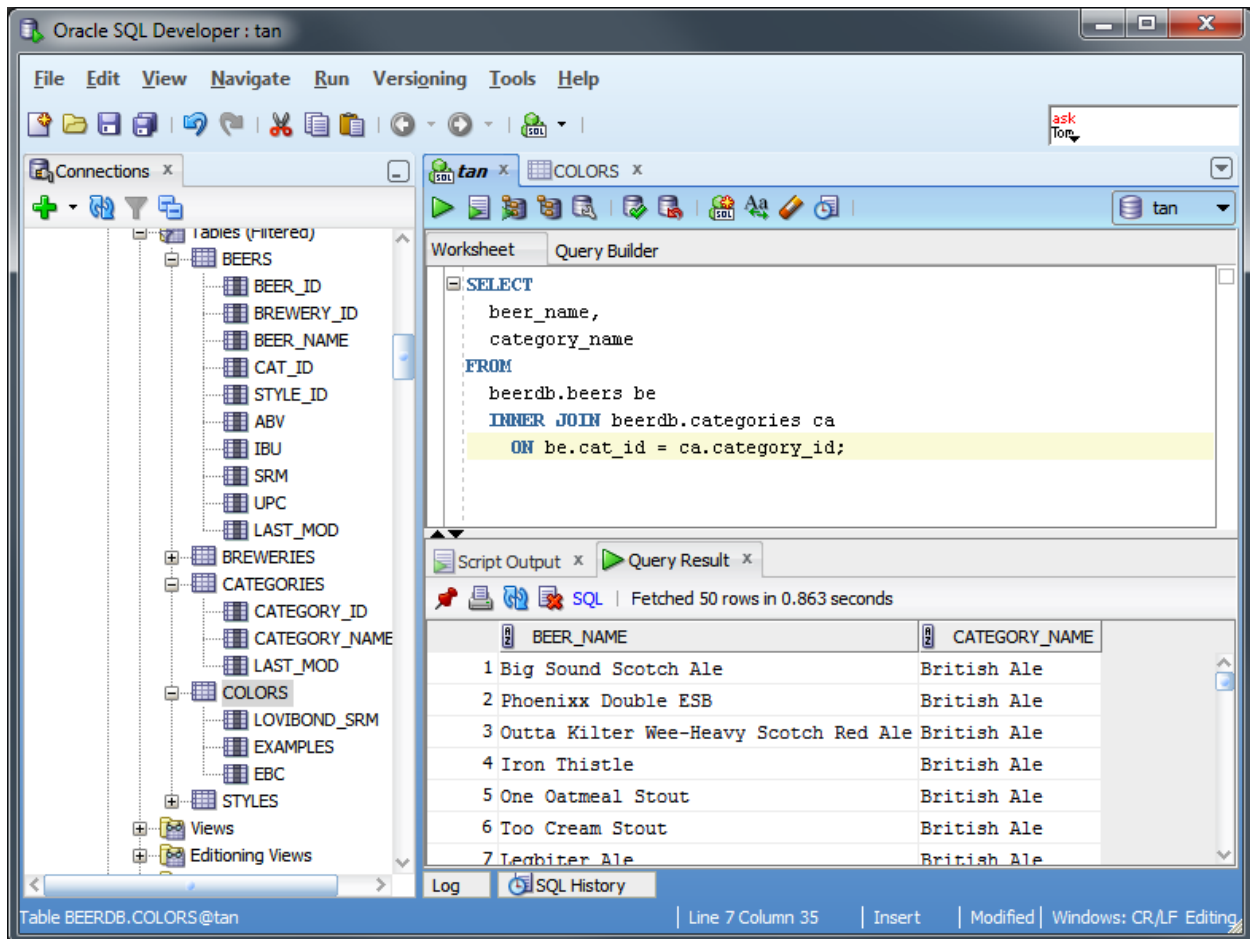
The 'Script Output' pane shows the query results, fetched in 0.04 seconds. The results are displayed in a table with two columns: BEER\_NAME and STYLE\_NAME. A blue arrow points from the BEER\_NAME column to the STYLE\_NAME column, highlighting the NULL values in the style names for the first two rows.

BEER_NAME	STYLE_NAME
1 Irish Red	(null)
2 Belgian Wit	(null)
3 EOS Hefeweizen	South German-Style P
4 India Pale Ale	American-Style Pale
5 Winter Ale	Old Ale
6 Pride & Joy Mild Ale	English-Style Pale

## Writing a Final Beer Database Query

We can tackle one final beer database query to demonstrate multiple joins. The query specification is as follows. Clearly, we need to combine BEERS with CATEGORIES and COLORS.

*“Display each beer’s name, category name, color example, and style name, for all beers that have values for category name, color example, and style name.”*



The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows a connection to 'tan'. The 'Tables (Filtered)' pane shows a tree view of the database schema, including tables like BEERS, CATEGORIES, and COLORS. The 'Query Builder' pane shows the following SQL query:

```
SELECT
  beer_name,
  category_name
FROM
  beerdbe.beers be
  INNER JOIN beerdbe.categories ca
    ON be.cat_id = ca.category_id;
```

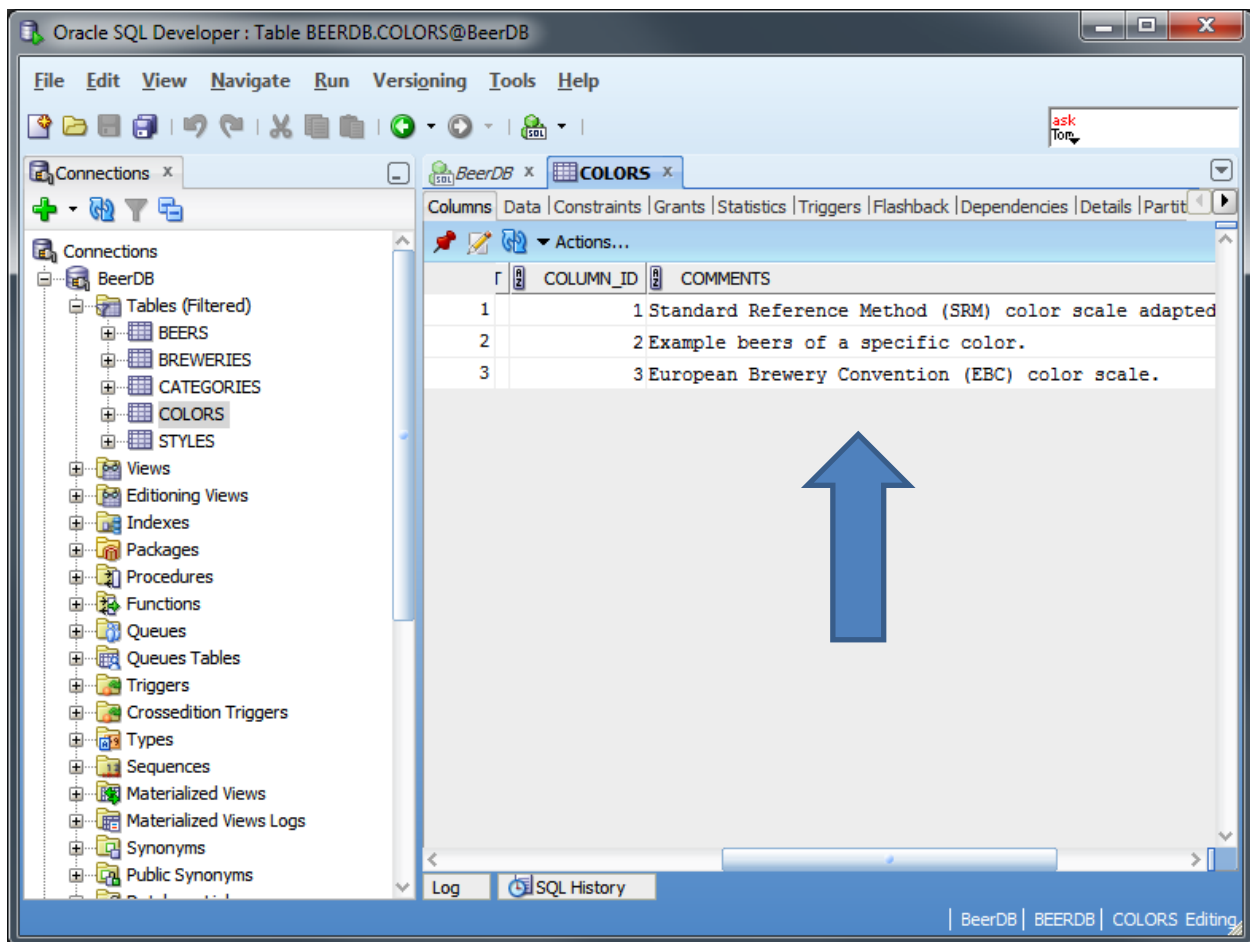
The 'Query Result' pane shows the results of the query, fetched in 0.863 seconds. The results are displayed in a table with two columns: BEER\_NAME and CATEGORY\_NAME.

BEER_NAME	CATEGORY_NAME
1 Big Sound Scotch Ale	British Ale
2 Phoenixx Double ESB	British Ale
3 Outta Kilter Wee-Heavy Scotch Red Ale	British Ale
4 Iron Thistle	British Ale
5 One Oatmeal Stout	British Ale
6 Too Cream Stout	British Ale
7 Leghitter Ale	British Ale

We start by joining the BEERS table with the CATEGORIES table using an INNER JOIN on CATEGORY\_ID. Note that the column names can differ, with CAT\_ID in the BEERS table and CATEGORY\_ID in the CATEGORIES table. That is fine as long as there really is a foreign key relationship between these tables and columns. Secondly, we again use table aliases (“be” and “ca”) to save some typing and make the query easier to interpret.

Next we have to add another join to handle the COLORS table. Remember that a join takes two relations and creates a single output relation. So, we take the output of the first join and pair that with another table (COLORS).

Before we proceed, we can explore the COLORS table a bit. Clicking on the COLORS table shows some details. If you scroll over on the “Columns” tab, you will see a COMMENTS column. This is metadata about the table. The Oracle DBMS provides a [COMMENT](#) statement for specifying table or column comments. This is a very nice way of documenting a database and helping future query writers. Here we can see that the SRM column refers to the [Standard Reference Method \(SRM\)](#) used to specify the color of beer. This also takes the mystery out of EBC, an alternate European Brewery Convention for beer color.



Now let's get back to writing the query. The next step is to add another join, bring the COLORS table into the mix. Again, we use an INNER JOIN.

This SQL code below shows the two joins (and three tables). Note that again the COLORS table is joined with columns that have slightly different names, SRM and LOVIBOND\_SRM. Now we know why from the column comments above. The EXAMPLES column shows the types of beers that typify this color. This is a helpful way of describing beer color. Not many beers in the database have color examples, but the query is kind of nice.

The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows a tree view of the database schema, including tables like BEERS, BEER\_ID, BREWERY\_ID, BEER\_NAME, CAT\_ID, STYLE\_ID, ABV, IBU, SRM, UPC, LAST\_MOD, BREWERIES, CATEGORIES, CATEGORY\_ID, CATEGORY\_NAME, LAST\_MOD, COLORS, LOVIBOND\_SRM, EXAMPLES, EBC, and STYLES. The 'Query Builder' pane on the right shows the following SQL query:

```
SELECT
  beer_name,
  category_name,
  examples
FROM
  beerdb.beers be
  INNER JOIN beerdb.categories ca
    ON be.cat_id = ca.category_id
  INNER JOIN beerdb.colors co
    ON be.srm = co.lovibond_srm;
```

Below the query, the 'Script Output' pane shows the results of the query. The status bar indicates 'All Rows Fetched: 6 in 0.047 seconds'. The results are displayed in a table with three columns: BEER\_NAME, CATEGORY\_NAME, and EXAMPLES.

	BEER_NAME	CATEGORY_NAME	EXAMPLES
1	Conundrum Red Ale	North American Ale	Biere de Garde, Double IPA
2	Christmas Ale	North American Ale	Foreign Stout, Baltic Porter
3	Dos Perros	North American Ale	Biere de Garde, Double IPA
4	Pale Ale	North American Ale	American Pale Ale, India Pal
5	Sly Rye Porter	North American Ale	Stout

That's it for now. Hopefully, this "cheat sheet" helps you complete the query writing assignment and learn a little SQL along the way. Please feel free to ask further questions or suggest additions to this document as you work on other queries. Good luck.